

[Artifact391] Programmable MCMC with Soundly Composed Guide Programs

Long Pham, Di Wang, Feras Saad, Jan Hoffmann

July 2, 2024

1 What is Included

From Zenodo, you should obtain the following items:

- This `README.pdf` document
- Archived and compressed Docker image `guide.types.tar.gz`

2 Getting Started Guide

Our artifact is a program analysis tool with the following capabilities:

1. Check structural type equality of (possibly context-free) guide types
2. Infer guide types of model and guide programs (using the structural-type-equality checking algorithm)
3. Check whether the support of sequentially composed guide programs coincides with the support of a model program.

In probabilistic programming with our new coroutine-based programmable inference framework, the user provides a model program and a sequential composition of guide programs. The model program specifies a probabilistic model for Bayesian inference. The sequential composition of guide programs customizes the Block Metropolis-Hastings (BMH) algorithm, where we successively run the guide programs, each of which is followed by an MH acceptance routine. Each guide program only updates a subset (i.e., block) of random variables. The model and guide programs are coroutines that communicate with other by message passing, and their communication protocols are described by guide types. To algorithmically decide structural type equality of guide types, our artifact implements the bisimilarity-checking algorithm by Hirshfeld and Moller [1] for context-free processes with finite norms, which model guide types. This structural-type-equality checking algorithm is also incorporated into the type-inference algorithm for guide types. Finally, to verify the soundness of the BMH algorithm, our artifact checks

whether the support of the sequentially composed guide programs covers all traces all possible traces admitted by the model program.

The artifact is wrapped inside the accompanying Docker image `guide_types.tar.gz`. Before running it, first install Docker as instructed here: <https://docs.docker.com/engine/install/>. To see if Docker has been installed properly, run

```
$ docker --version
Docker version 27.0.3, build 7d4bcd8
```

Load a Docker image by running

```
$ docker load --input guide_types.tar.gz
```

It creates an image named `guide_types` and stores it locally on your Docker may create an image with a slightly different name from `guide_types`. To check the name of the image, display all Docker images on your local machine by running

```
$ docker images
```

To run the image `guide_types`, run

```
$ docker run --name guide_types -it --rm guide_types
root@5b1c8c873064:/home/GuideTypes#
```

It creates a Docker container (i.e., a runnable instance of the Docker image) named `guide_types` and starts a shell inside the container. If the command does not run properly, you can instead build the image locally on your machine as instructed in ??.

Throughout this document, any command line starting with `#` is executed inside the Docker container, and any command line starting with `$` is executed in your local machine's terminal.

2.1 Structural-Type-Equality Checking

The initial working directory `type-equality` contains a file `type-equality-sample` storing several guide-type definitions. To display the content of the file, we run

```
# cat type-equality-sample
```

It prints out

```
type T1 = &{ $ | real /\ T1[T1] }
type T2 = &{ $ | real /\ T2[T2] }
type T3 = &{ $ | real /\ real /\ T3[T3] }
```

The first line defines a guide-type operator $T_1[\cdot]$ as

$$T_1[X] := X \& (\mathbb{R} \wedge \mathbb{R} \wedge T_1[T_1[X]]),$$

where X is a type variable that stands for a continuation type (i.e., the type of the communication protocol that we run after T_1 is finished). The type T_1 means the coroutine

receives a branch selection (from another coroutine that it communicates with), and proceeds to either X (i.e., the continuation guide type) or $\mathbb{R} \wedge T_1[T_1[X]]$, depending on which branch is taken. The guide type $\mathbb{R} \wedge T_1[T_1[X]]$ means the coroutine first sends a message of type \mathbb{R} (i.e., real numbers) and then proceeds to the guide type $T_1[T_1[X]]$. Here, $T_1[T_1[X]]$ can be interpreted as a sequential composition of two instances of T_1 , followed by the original continuation X .

To algorithmically decide the structural type equality between the guide-type operators $T_1[X]$ and $T_2[X]$, run

```
# dune exec gtypes type-equality type-equality-sample T1 T2
```

The command first prints out information about intermediate results of the type-equality checking algorithm for debugging. At the end of the printout, the command displays the result of structural-type-equality checking:

```
Types T1 and T2 are equal
```

This result is as expected because the two type operators $T_1[X]$ and $T_2[X]$ represent the same type—they only differ in the names of type operators.

Next, to check the structural type equality between $T_1[X]$ and $T_3[X]$, run

```
# dune exec gtypes type-equality type-equality-sample T1 T3
```

At the end of the printout, it displays:

```
Types T1 and T3 are unequal
```

This is a correct result because the right branch of $T_1[X]$ only sends one value (indicated by $\mathbb{R} \wedge$) before possible termination, while the right branch of $T_3[X]$ sends two values before possible termination.

2.2 Guide-Type Inference

Go to the directory `/home/GuideTypes/bench/coverage-check/recursive/recur` by running

```
# cd /home/GuideTypes/bench/coverage-check/recursive/recur
```

It contains a benchmark program `recur-covered-aligned`. To display its content, run

```
# cat recur-covered-aligned
```

The first three lines of the file are

```
type Old_trace
type General
type General_no_old
```

They declares three type operators that are later mentioned in the type signatures of coroutines. While the user needs to manually declare these type operators and insert them to the guide programs' type signatures, the type operators' definitions will be automatically inferred by our artifact.

The file defines three coroutines: `Recur1`, `Recur2`, and `Recur_no_old`. In each of them, the first line states the type signature of the coroutine. For example, the type signature of the the coroutine `Recur1` is

```
proc Recur1() -> unit | old : Old_trace | lat : General
```

It states (i) the coroutine `Recur1` has the output type of `unit`, (ii) the coroutine uses a channel named `old` of the guide type `Old_trace`, and (iii) the coroutine uses a channel named `lat` of the guide type `lat`.

To infer the definitions of the guide types mentioned in this file, we run

```
# dune exec gtypes type-check recur-covered-aligned
```

It prints out the inferred definitions of the three type operators. For instance, the inference result for the type operator `General` is stated as:

```
successfully inferred guide type operator General[$]
&{ $ | real /\ General[real /\ General[real /\ General[$]] } }
```

It translates to

$$\text{General}[X] := X \ \& \ (\mathbb{R} \wedge \text{General}[\mathbb{R} \wedge \text{General}[\mathbb{R} \wedge \text{General}[X]]]).$$

In addition to the inference results of guide types, the command also prints out (i) a list of pairs of guide types whose structural type equality we need to check during type inference and (ii) whether each pair is indeed equal. An example is

```
Types $ and $ are equal modulo coverage
```

At the end of the command's output, it displays `All equal modulo coverage`. It means all type-equality conditions discharged during type inference are verified, and hence the three inferred type definitions are also valid.

2.3 Coverage Checking

The last two liens in the benchmark program `recur-covered-aligned` define how the guide coroutines are sequentially composed:

```
Initial_type: Init_type
Guide_composition: Recur1; Recur2
```

The first line states the initial guide type `Init_type`, which is defined in the fourth line in the file:

```
type Init_type = &{ $ | real_u /\ Init_type[real_u /\
                    Init_type[real_u /\ Init_type[$]] } }
```

This guide type describes the initial trace fed to the sequential composition of guide coroutines. Because we assume the initial trace is uncovered (i.e., all random variables in the trace haven't been freshly sampled), the guide type `Init_type` contains `real_u`, where the subscript `u` means “uncovered.”

Finally, the last line in the file, `Guide_composition: Recur1; Recur2`, means we sequentially compose the guide coroutines `Recur1` and `Recur2` in this order.

To verify the full coverage of the sequential composition of the two guide coroutines, we run

```
# dune exec gtypes coverage-check recur-covered-aligned
```

The result is displayed at the end of the printout:

```
The final type is fully covered
```

References

- [1] Y. Hirshfeld and F. Moller. A Fast Algorithm for Deciding Bisimilarity of Normed Context-Free Processes. In B. Jonsson and J. Parrow, editors, *CONCUR '94: Concurrency Theory*, Lecture Notes in Computer Science, pages 48–63, Berlin, Heidelberg, 1994. Springer. ISBN 978-3-540-48654-1. doi: 10.1007/978-3-540-48654-1_5.

[1](#)